


```
!pip install -q kaggle
```

```
from google.colab import files
files.upload()
```

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.


Saving kaggle (2).json to kaggle (2).json

```
{'kaggle (2) json': {'username': 'cpazvhuo' 'key': 'a2h356a8d172h567d86528da0d43e8a6'}}}
```


```
!mkdir -p ~/.kaggle
```

```
!cp kaggle.json ~/.kaggle/
```

```
!chmod 600 ~/.kaggle/kaggle.json
```

 cp: cannot stat 'kaggle.json': No such file or directory
chmod: cannot access '/root/.kaggle/kaggle.json': No such file or directory

```
!kaggle datasets download -d rhammell/ships-in-satellite-imagery
```

 Dataset URL: <https://www.kaggle.com/datasets/rhammell/ships-in-satellite-imagery>
License(s): CC-BY-SA-4.0
Downloading ships-in-satellite-imagery.zip to /content
96% 178M/185M [00:01<00:00, 125MB/s]
100% 185M/185M [00:01<00:00, 130MB/s]

```
!unzip ships-in-satellite-imagery.zip
```



```

inflating: shipsnet/shipsnet/0_20161218_180844_0e26_-122.43818561434492_37.88189597797097.png
inflating: shipsnet/shipsnet/0_20161218_180844_0e26_-122.44292051246497_37.86017914204665.png
inflating: shipsnet/shipsnet/0_20161218_180844_0e26_-122.4796958662502_37.829549813389136.png
inflating: shipsnet/shipsnet/0_20161218_180844_0e26_-122.47984174489888_37.83790556670464.png
inflating: shipsnet/shipsnet/0_20161218_180844_0e26_-122.48871932039954_37.84589395485927.png

```

```
!ls
```

```

📁 'kaggle (2).json'  scenes          shipsnet
sample_data        ships-in-satellite-imagery.zip  shipsnet.json

```

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.utils import to_categorical

```

```
!ls /content
```

```

📁 'kaggle (2).json'  scenes          shipsnet
sample_data        ships-in-satellite-imagery.zip  shipsnet.json

```

```
!unzip /content/ships-in-satellite-imagery.zip -d /content/ships-in-satellite-imagery
```

```
📁
```

```

inflating: /content/ships-in-satellite-imagery/shipsnet/shipsnet/0__20161102_180658_0e26__-122.23122736968294_37.75239959943399.png
inflating: /content/ships-in-satellite-imagery/shipsnet/shipsnet/0__20161102_180658_0e26__-122.28126318568829_37.769020038003475.png
inflating: /content/ships-in-satellite-imagery/shipsnet/shipsnet/0__20161102_180658_0e26__-122.31527680074582_37.79277375149748.png
inflating: /content/ships-in-satellite-imagery/shipsnet/shipsnet/0__20161102_180658_0e26__-122.32843644745485_37.73923054907409.png
inflating: /content/ships-in-satellite-imagery/shipsnet/shipsnet/0__20161102_180658_0e26__-122.34040514248915_37.748761931300486.png

```

```
!ls /content/ships-in-satellite-imagery
```

```
📁 scenes shipsnet shipsnet.json
```

```
dataset_path = "/content/ships-in-satellite-imagery/shipsnet"
```

```

import os
import numpy as np
from tensorflow.keras.preprocessing.image import load_img, img_to_array
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

# Path to the dataset (update this if necessary)
dataset_path = "/content/ships-in-satellite-imagery"

# Verify the dataset path
if not os.path.exists(dataset_path):
    # Debug: Print available directories
    print("Available directories in /content:")
    print(os.listdir("/content"))
    raise FileNotFoundError(f"The dataset path '{dataset_path}' does not exist. Please check the path.")

# Debug: Print contents of the dataset directory
print("Contents of the dataset directory:")
print(os.listdir(dataset_path))

# Load images and labels
images = []
labels = []

# Recursively search for .png files
for root, dirs, files in os.walk(dataset_path):
    for filename in files:
        if filename.endswith(".png"): # Ensure only PNG files are processed
            # Load image
            img_path = os.path.join(root, filename)
            img = load_img(img_path, target_size=(80, 80)) # Resize images to 80x80
            img = img_to_array(img) / 255.0 # Normalize pixel values to [0, 1]
            images.append(img)

            # Extract label from filename (e.g., "ship_1.png" -> "ship")
            label = filename.split("_")[0] # Adjust based on filename structure
            labels.append(label)

# Check if any images were loaded
if len(images) == 0:
    raise ValueError("No images were loaded. Check the dataset path and file structure.")

# Convert to numpy arrays
images = np.array(images)
labels = np.array(labels)

# Encode labels to integers
label_encoder = LabelEncoder()
labels_encoded = label_encoder.fit_transform(labels)

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(images, labels_encoded, test_size=0.2, random_state=42)

print(f"Loaded {len(images)} images and {len(labels)} labels.")
print(f"Unique labels: {np.unique(labels)}")
print(f"Training data shape: {X_train.shape}, Testing data shape: {X_test.shape}")

```

```

📁 Contents of the dataset directory:
['shipsnet.json', 'shipsnet', 'scenes']
Loaded 4008 images and 4008 labels.
Unique labels: ['0' '1' 'lb' 'sfbay']
Training data shape: (3206, 80, 80, 3), Testing data shape: (802, 80, 80, 3)

```

```

import os
import numpy as np
from tensorflow.keras.preprocessing.image import load_img, img_to_array
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

dataset_path = "/content/ships-in-satellite-imagery"
if not os.path.exists(dataset_path):
    raise FileNotFoundError(f"The dataset path '{dataset_path}' does not exist. Please check the path.")

images = []
labels = []

for filename in os.listdir(dataset_path):
    if filename.endswith(".png"): # Ensure only PNG files are processed
        # Load image
        img_path = os.path.join(dataset_path, filename)
        img = load_img(img_path, target_size=(80, 80)) # Resize images to 80x80
        img = img_to_array(img) / 255.0 # Normalize pixel values to [0, 1]
        images.append(img)

        # Extract label from filename (e.g., "ship_1.png" -> "ship")
        label = filename.split("_")[0] # Adjust based on filename structure
        labels.append(label)

images = np.array(images)
labels = np.array(labels)

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

# Define the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(80, 80, 3)), # First convolutional layer
    MaxPooling2D((2, 2)), # First max-pooling layer
    Conv2D(64, (3, 3), activation='relu'), # Second convolutional layer
    MaxPooling2D((2, 2)), # Second max-pooling layer
    Conv2D(128, (3, 3), activation='relu'), # Third convolutional layer
    MaxPooling2D((2, 2)), # Third max-pooling layer
    Flatten(), # Flatten the output for dense layers
    Dense(128, activation='relu'), # Fully connected layer
    Dropout(0.5), # Dropout to prevent overfitting
    Dense(len(label_encoder.classes_), activation='softmax') # Output layer
])

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Print model summary
model.summary()

```

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` to `input` in the constructor of `Conv2D` or `Conv3D` layers. Instead, use the `input_shape` argument in the `call` method.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 78, 78, 32)	896
max_pooling2d (MaxPooling2D)	(None, 39, 39, 32)	0
conv2d_1 (Conv2D)	(None, 37, 37, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 18, 18, 64)	0
conv2d_2 (Conv2D)	(None, 16, 16, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 128)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 128)	1,048,704
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 4)	516

Total params: 1,142,468 (4.36 MB)

Trainable params: 1,142,468 (4.36 MB)

```
from tensorflow.keras.callbacks import EarlyStopping
```

```
# Data Augmentation
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
datagen = ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)
```

```
# Fit the data generator to the training data
datagen.fit(X_train)
```

```
# Early stopping to prevent overfitting
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
```

```
# Train the model
history = model.fit(
    datagen.flow(X_train, y_train, batch_size=32),
    epochs=50,
    validation_data=(X_test, y_test),
    callbacks=[early_stopping]
)
```

```

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should inherit from `PyDatasetAdapter` and implement the `get_data_adapter_class` method. self.warn if super not called()

```

Epoch 1/50	101/101	63s	596ms/step	- accuracy: 0.7100	- loss: 0.6549	- val_accuracy: 0.8890	- val_loss: 0.3661
Epoch 2/50	101/101	57s	556ms/step	- accuracy: 0.8127	- loss: 0.4288	- val_accuracy: 0.9015	- val_loss: 0.3399
Epoch 3/50	101/101	60s	591ms/step	- accuracy: 0.8301	- loss: 0.4136	- val_accuracy: 0.8878	- val_loss: 0.3486
Epoch 4/50	101/101	60s	599ms/step	- accuracy: 0.8408	- loss: 0.3879	- val_accuracy: 0.9052	- val_loss: 0.3225
Epoch 5/50	101/101	75s	527ms/step	- accuracy: 0.8452	- loss: 0.3505	- val_accuracy: 0.8791	- val_loss: 0.3148
Epoch 6/50	101/101	59s	588ms/step	- accuracy: 0.8443	- loss: 0.3523	- val_accuracy: 0.8940	- val_loss: 0.3188
Epoch 7/50	101/101	61s	605ms/step	- accuracy: 0.8412	- loss: 0.3459	- val_accuracy: 0.8853	- val_loss: 0.3211
Epoch 8/50	101/101	55s	543ms/step	- accuracy: 0.8463	- loss: 0.3343	- val_accuracy: 0.8990	- val_loss: 0.3080
Epoch 9/50	101/101	81s	538ms/step	- accuracy: 0.8575	- loss: 0.3300	- val_accuracy: 0.8853	- val_loss: 0.3085
Epoch 10/50	101/101	58s	578ms/step	- accuracy: 0.8574	- loss: 0.3141	- val_accuracy: 0.8990	- val_loss: 0.2949

```

Epoch 11/50
101/101 ----- 53s 524ms/step - accuracy: 0.8610 - loss: 0.3108 - val_accuracy: 0.9077 - val_loss: 0.2806
Epoch 12/50
101/101 ----- 59s 585ms/step - accuracy: 0.8614 - loss: 0.2963 - val_accuracy: 0.9152 - val_loss: 0.2631
Epoch 13/50
101/101 ----- 53s 529ms/step - accuracy: 0.8816 - loss: 0.2730 - val_accuracy: 0.8778 - val_loss: 0.3068
Epoch 14/50
101/101 ----- 62s 612ms/step - accuracy: 0.8861 - loss: 0.2747 - val_accuracy: 0.9314 - val_loss: 0.2335
Epoch 15/50
101/101 ----- 57s 566ms/step - accuracy: 0.8913 - loss: 0.2592 - val_accuracy: 0.9314 - val_loss: 0.2207
Epoch 16/50
101/101 ----- 57s 564ms/step - accuracy: 0.8921 - loss: 0.2490 - val_accuracy: 0.9364 - val_loss: 0.1966
Epoch 17/50
101/101 ----- 62s 619ms/step - accuracy: 0.8986 - loss: 0.2350 - val_accuracy: 0.9289 - val_loss: 0.2175
Epoch 18/50
101/101 ----- 78s 577ms/step - accuracy: 0.8833 - loss: 0.2706 - val_accuracy: 0.9339 - val_loss: 0.2110
Epoch 19/50
101/101 ----- 55s 549ms/step - accuracy: 0.8962 - loss: 0.2408 - val_accuracy: 0.9327 - val_loss: 0.1968
Epoch 20/50
101/101 ----- 57s 557ms/step - accuracy: 0.8983 - loss: 0.2476 - val_accuracy: 0.9439 - val_loss: 0.1793
Epoch 21/50
101/101 ----- 57s 563ms/step - accuracy: 0.9123 - loss: 0.2022 - val_accuracy: 0.9439 - val_loss: 0.1733
Epoch 22/50
101/101 ----- 56s 551ms/step - accuracy: 0.9016 - loss: 0.2376 - val_accuracy: 0.9214 - val_loss: 0.2247
Epoch 23/50
101/101 ----- 56s 558ms/step - accuracy: 0.9090 - loss: 0.2281 - val_accuracy: 0.9414 - val_loss: 0.1686
Epoch 24/50
101/101 ----- 56s 560ms/step - accuracy: 0.9040 - loss: 0.2263 - val_accuracy: 0.9476 - val_loss: 0.1676
Epoch 25/50
101/101 ----- 57s 566ms/step - accuracy: 0.9264 - loss: 0.2086 - val_accuracy: 0.9464 - val_loss: 0.1543
Epoch 26/50
101/101 ----- 55s 545ms/step - accuracy: 0.9128 - loss: 0.2122 - val_accuracy: 0.9426 - val_loss: 0.1479
Epoch 27/50
101/101 ----- 56s 556ms/step - accuracy: 0.9081 - loss: 0.2285 - val_accuracy: 0.9464 - val_loss: 0.1544
Epoch 28/50
101/101 ----- 56s 556ms/step - accuracy: 0.9081 - loss: 0.2285 - val_accuracy: 0.9464 - val_loss: 0.1544

```

```

# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")

# Plot training and validation accuracy and loss
import matplotlib.pyplot as plt

plt.figure(figsize=(12, 5))



# Plot accuracy
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

# Plot loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.show()

```

26/26 ————— 2s 89ms/step - accuracy: 0.9565 - loss: 0.1425
 Test Accuracy: 95.14%



```
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import GlobalAveragePooling2D

# Load pre-trained VGG16 model (without the top layers)
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(80, 80, 3))

# Freeze the base model layers
base_model.trainable = False


# Add custom layers on top of the base model
model = Sequential([
    base_model,
    GlobalAveragePooling2D(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(len(label_encoder.classes_), activation='softmax')
])

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(
    datagen.flow(X_train, y_train, batch_size=32),
    epochs=20,
    validation_data=(X_test, y_test),
    callbacks=[early_stopping]
)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.58889256/58889256 0s 0us/step

```
Epoch 1/20
101/101 ————— 397s 4s/step - accuracy: 0.7243 - loss: 0.6724 - val_accuracy: 0.9152 - val_loss: 0.2355
Epoch 2/20
101/101 ————— 393s 4s/step - accuracy: 0.8917 - loss: 0.2642 - val_accuracy: 0.9352 - val_loss: 0.1930
Epoch 3/20
101/101 ————— 391s 4s/step - accuracy: 0.9156 - loss: 0.2034 - val_accuracy: 0.9426 - val_loss: 0.1740
Epoch 4/20
101/101 ————— 394s 4s/step - accuracy: 0.9190 - loss: 0.2010 - val_accuracy: 0.9352 - val_loss: 0.1874
Epoch 5/20
101/101 ————— 394s 4s/step - accuracy: 0.9256 - loss: 0.1849 - val_accuracy: 0.9401 - val_loss: 0.1757
Epoch 6/20
101/101 ————— 392s 4s/step - accuracy: 0.9407 - loss: 0.1592 - val_accuracy: 0.9514 - val_loss: 0.1464
Epoch 7/20
101/101 ————— 386s 4s/step - accuracy: 0.9389 - loss: 0.1502 - val_accuracy: 0.9564 - val_loss: 0.1387
Epoch 8/20
101/101 ————— 395s 4s/step - accuracy: 0.9320 - loss: 0.1689 - val_accuracy: 0.9489 - val_loss: 0.1484
Epoch 9/20
101/101 ————— 386s 4s/step - accuracy: 0.9414 - loss: 0.1560 - val_accuracy: 0.9551 - val_loss: 0.1465
Epoch 10/20
```

101/101  394s 4s/step - accuracy: 0.9355 - loss: 0.1532 - val_accuracy: 0.9626 - val_loss: 0.1168